CLAUDE OPUS 4.7

Same sticker, Bigger bill.

Claude Opus 4.7 raised per-session cost roughly 20 to 30% on measured workloads. The per-token price did not change.

SAME TEXT. DIFFERENT COUNTS.

`fetchUserData`

CLAUDE 4.6	CLAUDE 4.7
fetch User Data	fet ch Use r Data
3 tokens	5 tokens

1.325x on real workloads

weighted average on authentic Claude Code traffic

Swipe to see the measurement →

MILIND NAIR

WHAT CHANGED

Same words. More tokens.

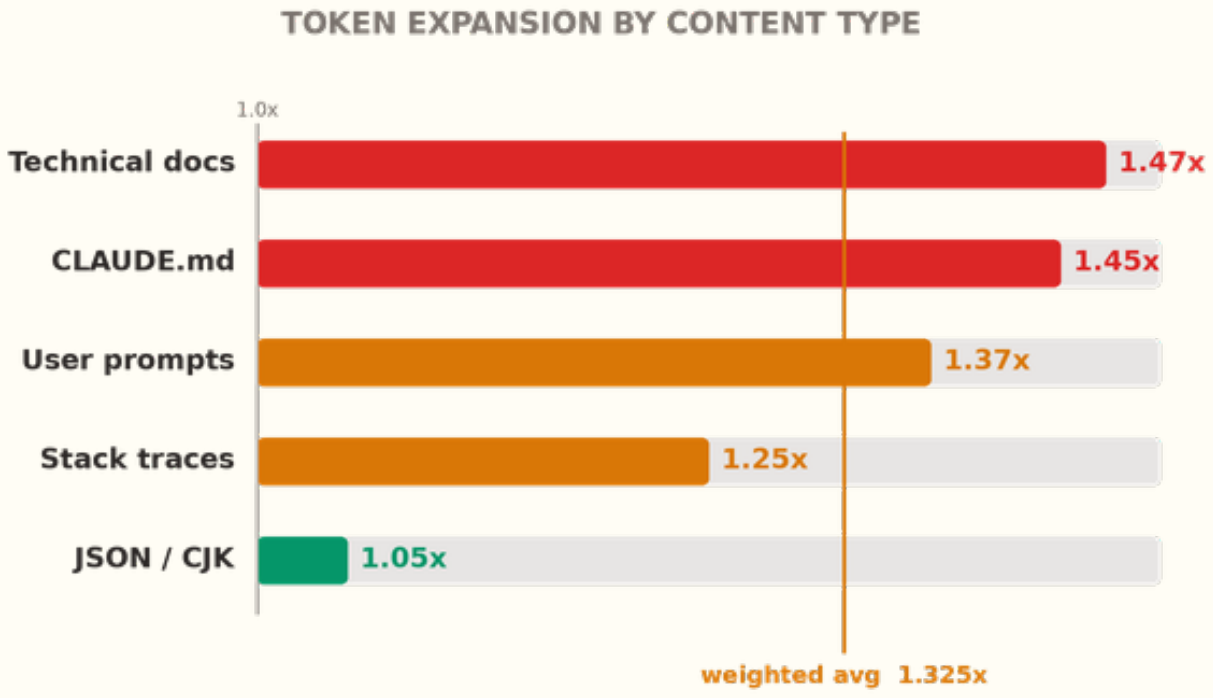

Measured on authentic Claude Code workloads

MEASURED AVERAGE

1.325x more tokens on authentic Claude Code workloads. CLAUDE.md files hit 1.45x.

THE BILL

Per-session cost rose roughly 20 to 30%.

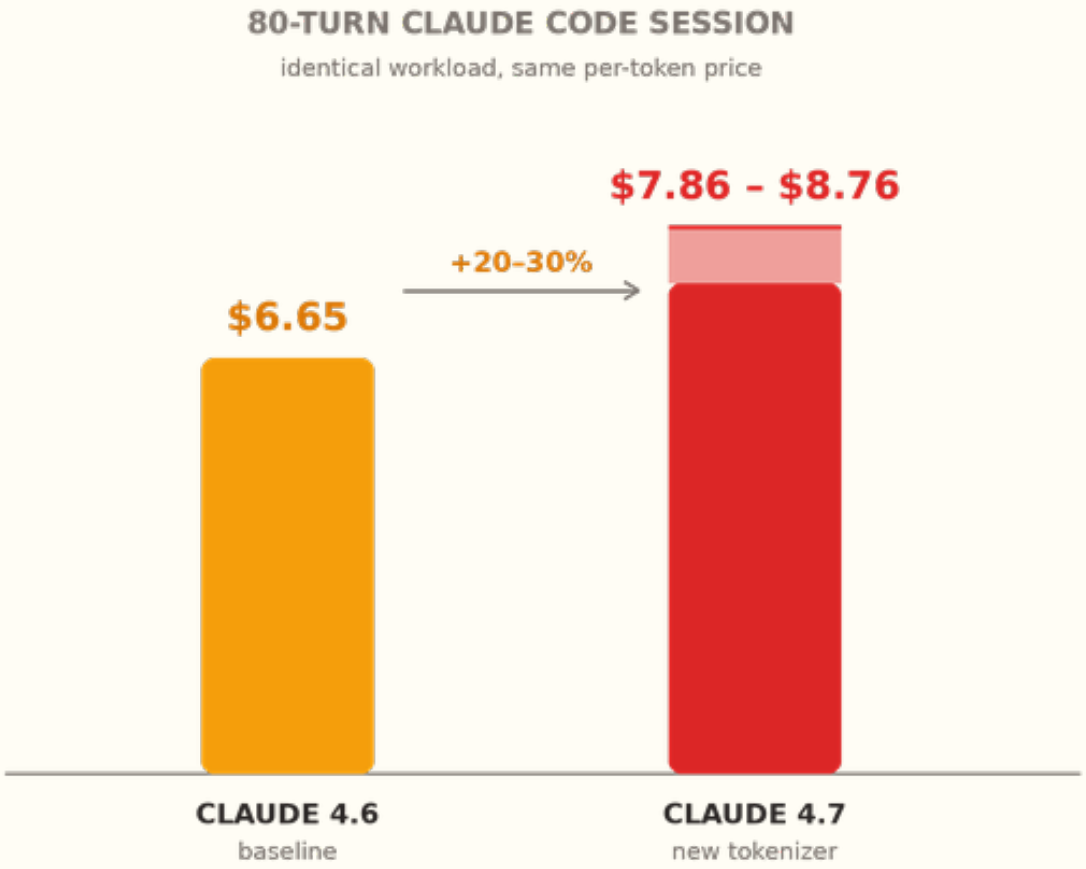

SAME WORK, BIGGER BILL

Per-token price held flat. Effective session cost climbed because the new tokenizer emitted more tokens per turn.

THE OTHER HALF

On IFEval, the gain was narrow.

IFEval, strict evaluation

Instruction-following improved 5 points, 17/20 to 18/20. N=20, so a one-prompt swing — directionally right, imprecise.

IFEval, loose evaluation

No change. Both Opus 4.6 and Opus 4.7 scored 18/20.

SCOPE NOTE

Opus 4.7's bigger wins sit on SWE-bench Pro (+10.9pp), MCP-Atlas (+14.6pp), and CursorBench (58 to 70), outside the tokenizer study's scope.

When the upgrade earns its keep.

1. **When is Opus 4.7 worth the cost uplift?**

   → When strict instruction-following matters, and the 5-point gain moves your pass rate across a threshold you care about.

2. **Where does the cost uplift hit hardest?**

   → On real-world samples: CLAUDE.md files expanded 1.45x, user prompts 1.37x, stack traces 1.25x. Prose-heavy context feels the pricing change most.

3. **What should you actually measure?**

   → Dollars per completed task. Per-token price is the sticker; session cost is the bill.

RULE OF THUMB

Tokens are the sticker. Tasks are the bill.

MILIND NAIR

Tokens are the sticker. Tasks are the bill.

MILIND NAIR